MSR113

Where NOT to Use AI in Health Economic Modeling: An Agentic Platform with Clear Trust Boundaries

Liza Zasukhina[1], Elena Meshkova[1], Illia Prazdnyk[1], Haidong Feng[2]

[1]MSD, Czech Republic; [2]Merck & Co., Inc., Boston, MA, USA

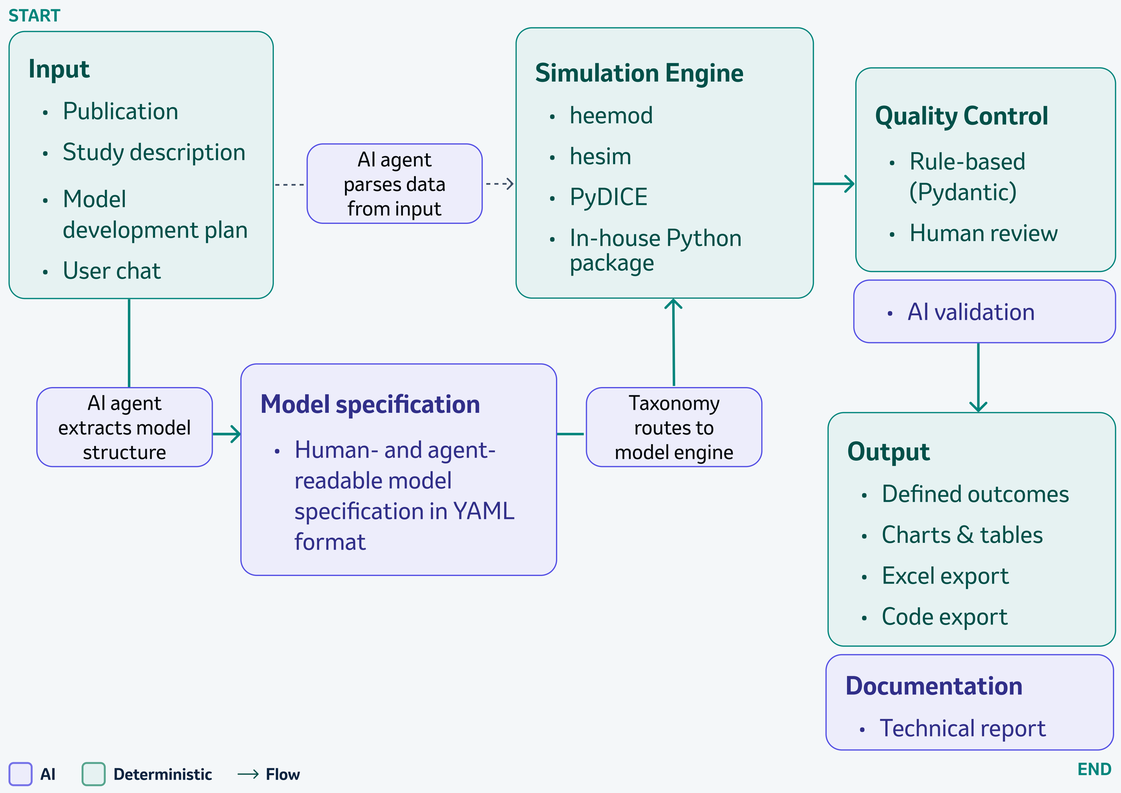
Background & Aims

- Health economists spend considerable time on evidence gathering, iterating model structures, and writing documentation — manual, repetitive tasks well suited to AI.
- However, AI-generated model logic, calculations, and even results are non-deterministic, prone to silent errors, and hard to QC, while HTA bodies expect transparency, reproducibility, and human accountability.
- We developed an AI workflow that accelerates the work around the model, but **never uses AI for computation**: simulation stays in established, deterministic packages.

Methods & Results

- Agents use tools such as Docling, which preserves table structure, to extract clinical and economic inputs from publications and HTA reports.
- Extracted inputs and model structure are captured in a human-readable YAML specification that modelers review and approve before simulation. It doubles as documentation and executable input, so changing the model structure (e.g., adding a health state) means editing the file, not recoding.

Fig. 1. AI workflow: AI structures inputs, deterministic engines compute.

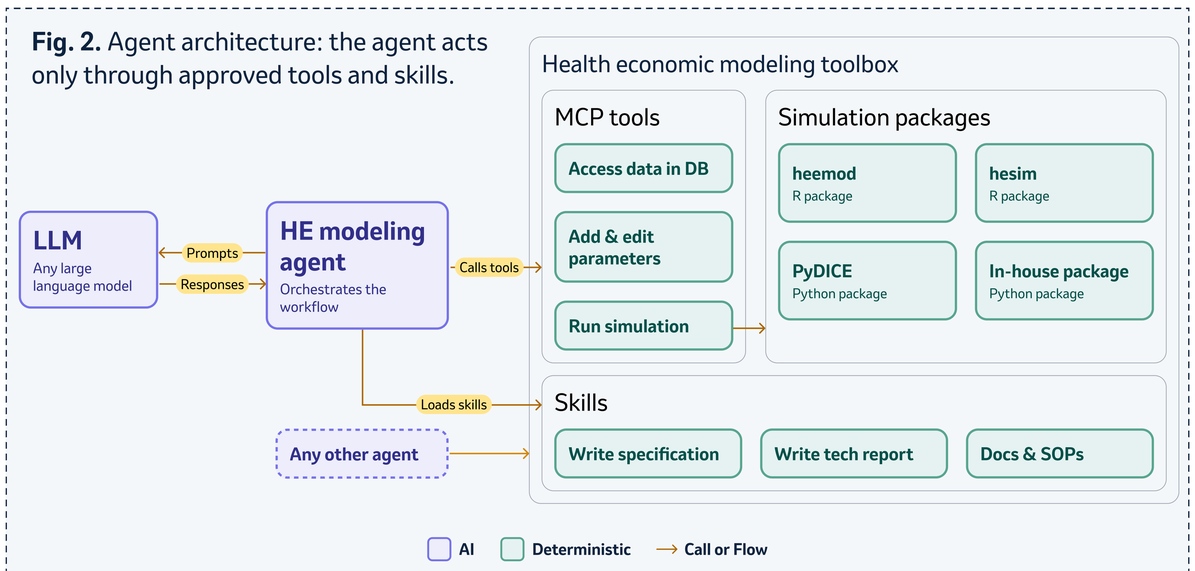


- Simulation is handled by established open-source packages — *heemod* and *hesim* in R, *PyDICE* in Python — and an in-house Python package, routed automatically by model taxonomy. The workflow can export executable code (R/Python) or Excel workbooks. All AI-extracted parameters remain traceable to source documents, and results are reproducible through the underlying packages (Fig. 1).

Methods & Results (cont.)

- In a proof-of-concept, the workflow went from a published study or HTA report to a running, reviewable simulation without manual coding — a step that typically takes weeks.
- Specifications are routed by model taxonomy to a suitable engine, e.g., cohort state-transition models to *heemod*, individual-level state-transition models to *hesim*, DICE models to *PyDICE*.
- Generated documentation maintained traceability to source evidence.
- The agent works inside a harness: it can act only through approved, quality-controlled tools and skills, exposed via MCP (Model Context Protocol). It calls the engines but **never writes or runs its own simulation code** (Fig. 2).

Fig. 2. Agent architecture: the agent acts only through approved tools and skills.



Implications & Conclusions

- We recommend using AI for the work around the model — evidence extraction, structuring, specification, documentation — and **deterministic tools for calculations and simulation code.** The same specification always gives the same results.
- Because every run traces back to a reviewed specification and a QC'd engine, health economists can shift focus from coding to what matters most: model structure and inputs.

Limitations & Next Steps

- Proof-of-concept scope: cost-effectiveness models only; not yet tested in an HTA submission.
- AI-drafted specifications may differ between runs or contain extraction errors, so human review remains essential before simulation.
- Next, an evidence agent continuously searching for new data could plug into the same system — a stepping stone toward living models.

Contact Information

Liza Zasukhina — liza.zasukhina@msd.com

Funding & disclosures: All authors are employees of MSD or Merck Sharp & Dohme LLC, Rahway, NJ, USA

References

1. Filipović-Pierucci A, et al. The R package heemod. arXiv:1702.03252; 2017.
2. Incerti D, Jansen JP. hesim: health economic simulation modeling. arXiv:2102.09437; 2021.
3. Caro JJ. DICE simulation for pharmacoeconomics. PharmacoEconomics 2016;34:665-72.
4. Auer C, et al. Docling technical report. arXiv:2408.09869; 2024.
5. NICE. Use of AI in evidence generation: NICE position statement. 2024.

Copyright © 2026 Merck & Co., Inc., Rahway, NJ, USA and its affiliates. All rights reserved.